



IBM Developer SKILLS NETWORK

Data Visualization

Estimated time needed: **30** minutes

Objectives

After completing this lab you will be able to:

- Create Data Visualization with Python
- Use various Python libraries for visualization

Introduction

The aim of these labs is to introduce you to data visualization with Python as concrete and as consistent as possible. Speaking of consistency, because there is no *best* data visualization library available for Python - up to creating these labs - we have to introduce different libraries and show their benefits when we are discussing new visualization concepts. Doing so, we hope to make students well-rounded with visualization libraries and concepts so that they are able to judge and decide on the best visualization technique and tool for a given problem *and* audience.

Please make sure that you have completed the prerequisites for this course, namely **Python Basics for Data Science** (https://www.edx.org/course/python-basics-for-data-science-2?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01) and **Analyzing Data with Python** (https://www.edx.org/course/data-analysis-with-python?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01).

Note: The majority of the plots and visualizations will be generated using data stored in *pandas* dataframes. Therefore, in this lab, we provide a brief crash course on *pandas*. However, if you are interested in learning more about the *pandas* library, detailed description and explanation of how to use it and how to clean, munge, and process data stored in a *pandas* dataframe are provided in our course **Analyzing Data with Python** (https://www.edx.org/course/data-analysis-with-python?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01).

Table of Contents

1. [Exploring Datasets with *pandas*](#0)
- 1.1 [The Dataset: Immigration to Canada from 1980 to 2013](#2)
- 1.2 [*pandas* Basics](#4)
- 1.3 [*pandas* Intermediate: Indexing and Selection](#6)
- 2\ [Visualizing Data using Matplotlib](#8)
- 2.1 [Matplotlib: Standard Python Visualization Library](#10)
- 3\ [Line Plots](#12)

Exploring Datasets with *pandas*

pandas is an essential data analysis toolkit for Python. From their [website](http://pandas.pydata.org/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01) (http://pandas.pydata.org/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01):

pandas is a Python package providing fast, flexible, and expressive data structures designed to make working with “relational” or “labeled” data both easy and intuitive. It aims to be the fundamental high-level building block for doing practical, **real world** data analysis in Python.

The course heavily relies on *pandas* for data wrangling, analysis, and visualization. We encourage you to spend some time and familiarize yourself with the *pandas* API Reference: <http://pandas.pydata.org/pandas-docs/stable/api.html> (http://pandas.pydata.org/pandas-docs/stable/api.html?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01).

The Dataset: Immigration to Canada from 1980 to 2013

```
import numpy as np # useful for many scientific computing in Python
import pandas as pd # primary data structure library
```

Let's download and import our primary Canadian Immigration dataset using *pandas*'s `read_excel()` method. Normally, before we can do that, we would need to download a module which *pandas* requires reading in Excel files. This module was **openpyxl** (formerly **xlrd**). For your convenience, we have pre-installed this module, so you would not have to worry about that. Otherwise, you would need to run the following line of code to install the **openpyxl** module:

```
! pip3 install openpyxl
```

Now we are ready to read in our data.

In [2]:

```
df_can = pd.read_excel(
    'https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DV
0101EN-SkillsNetwork/Data%20Files/Canada.xlsx',
    sheet_name='Canada by Citizenship',
    skiprows=range(20),
    skipfooter=2)

print('Data read into a pandas dataframe!')
```

Data read into a pandas dataframe!

Let's view the top 5 rows of the dataset using the `head()` function.

In [3]:

```
df_can.head()
# tip: You can specify the number of rows you'd like to see as follows: df_can.head(10)
```

Out[3]:

	Type	Coverage	OdName	AREA	AreaName	REG	RegName	DEV	DevName	1980	...	2004
0	Immigrants	Foreigners	Afghanistan	935	Asia	5501	Southern Asia	902	Developing regions	16	...	2978
1	Immigrants	Foreigners	Albania	908	Europe	925	Southern Europe	901	Developed regions	1	...	1450
2	Immigrants	Foreigners	Algeria	903	Africa	912	Northern Africa	902	Developing regions	80	...	3616
3	Immigrants	Foreigners	American Samoa	909	Oceania	957	Polynesia	902	Developing regions	0	...	0
4	Immigrants	Foreigners	Andorra	908	Europe	925	Southern Europe	901	Developed regions	0	...	0

5 rows × 43 columns



We can also view the bottom 5 rows of the dataset using the `tail()` function.


In [4]:

```
df_can.tail()
```

Out[4]:

	Type	Coverage	OdName	AREA	AreaName	REG	RegName	DEV	DevName	1980	...	2004
190	Immigrants	Foreigners	Viet Nam	935	Asia	920	South-Eastern Asia	902	Developing regions	1191	...	1816
191	Immigrants	Foreigners	Western Sahara	903	Africa	912	Northern Africa	902	Developing regions	0	...	0
192	Immigrants	Foreigners	Yemen	935	Asia	922	Western Asia	902	Developing regions	1	...	124
193	Immigrants	Foreigners	Zambia	903	Africa	910	Eastern Africa	902	Developing regions	11	...	56
194	Immigrants	Foreigners	Zimbabwe	903	Africa	910	Eastern Africa	902	Developing regions	72	...	1450

5 rows × 43 columns



When analyzing a dataset, it's always a good idea to start by getting basic information about your dataframe. We can do this by using the `info()` method.

This method can be used to get a short summary of the dataframe.

In [5]:

```
df_can.info(verbose=False)
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 195 entries, 0 to 194  
Columns: 43 entries, Type to 2013  
dtypes: int64(37), object(6)  
memory usage: 65.6+ KB
```

To get the list of column headers we can call upon the data frame's `columns` instance variable.

In [6]:

```
df_can.columns
```

Out[6]:

```
Index(['Type', 'Coverage', 'OdName', 'AREA', 'AreaName', 'REG',  
      'RegName', 'DEV', 'DevName', 1980, 1981, 1982,  
      1983, 1984, 1985, 1986, 1987, 1988,  
      1989, 1990, 1991, 1992, 1993, 1994,  
      1995, 1996, 1997, 1998, 1999, 2000,  
      2001, 2002, 2003, 2004, 2005, 2006,  
      2007, 2008, 2009, 2010, 2011, 2012,  
      2013],  
      dtype='object')
```

Similarly, to get the list of indices we use the `.index` instance variables.

In [7]:

```
df_can.index
```

Out[7]:

```
RangeIndex(start=0, stop=195, step=1)
```

Note: The default type of instance variables `index` and `columns` are **NOT** list .

In [8]:

```
print(type(df_can.columns))
print(type(df_can.index))
```

```
<class 'pandas.core.indexes.base.Index'>
```

```
<class 'pandas.core.indexes.range.RangeIndex'>
```

To get the index and columns as lists, we can use the `tolist()` method.

In [9]:

```
df_can.columns.tolist()
```

Out[9]:

```
['Type',  
 'Coverage',  
 'OdName',  
 'AREA',  
 'AreaName',  
 'REG',  
 'RegName',  
 'DEV',  
 'DevName',  
 1980,  
 1981,  
 1982,  
 1983,  
 1984,  
 1985,  
 1986,  
 1987,  
 1988,  
 1989,  
 1990,  
 1991,  
 1992,  
 1993,  
 1994,  
 1995,  
 1996,  
 1997,  
 1998,  
 1999,  
 2000,  
 2001,  
 2002,  
 2003,  
 2004,  
 2005,  
 2006,  
 2007,  
 2008,  
 2009,  
 2010,  
 2011,  
 2012,  
 2013]
```


In [10]:

```
df_can.index.tolist()
```

Out[10]:

```
[0,  
1,  
2,  
3,  
4,  
5,  
6,  
7,  
8,  
9,  
10,  
11,  
12,  
13,  
14,  
15,  
16,  
17,  
18,  
19,  
20,  
21,  
22,  
23,  
24,  
25,  
26,  
27,  
28,  
29,  
30,  
31,  
32,  
33,  
34,  
35,  
36,  
37,  
38,  
39,  
40,  
41,  
42,  
43,  
44,  
45,  
46,  
47,  
48,  
49,  
50,  
51,  
52,  
53,  
54,  
55,  
56,  
57,  
58,  
59,
```

60,
61,
62,
63,
64,
65,
66,
67,
68,
69,
70,
71,
72,
73,
74,
75,
76,
77,
78,
79,
80,
81,
82,
83,
84,
85,
86,
87,
88,
89,
90,
91,
92,
93,
94,
95,
96,
97,
98,
99,
100,
101,
102,
103,
104,
105,
106,
107,
108,
109,
110,
111,
112,
113,
114,
115,
116,
117,
118,
119,
120,
121,

122,
123,
124,
125,
126,
127,
128,
129,
130,
131,
132,
133,
134,
135,
136,
137,
138,
139,
140,
141,
142,
143,
144,
145,
146,
147,
148,
149,
150,
151,
152,
153,
154,
155,
156,
157,
158,
159,
160,
161,
162,
163,
164,
165,
166,
167,
168,
169,
170,
171,
172,
173,
174,
175,
176,
177,
178,
179,
180,
181,
182,
183,

184,
185,
186,
187,
188,
189,
190,
191,
192,
193,
194]

In [11]:

```
print(type(df_can.columns.tolist()))  
print(type(df_can.index.tolist()))
```

<class 'list'>
<class 'list'>

To view the dimensions of the dataframe, we use the `shape` instance variable of it.

In [12]:

```
# size of dataframe (rows, columns)  
df_can.shape
```

Out[12]:

(195, 43)

Note: The main types stored in *pandas* objects are `float` , `int` , `bool` , `datetime64[ns]` , `datetime64[ns, tz]` , `timedelta[ns]` , `category` , and `object` (string). In addition, these dtypes have item sizes, e.g. `int64` and `int32` .

Let's clean the data set to remove a few unnecessary columns. We can use *pandas* `drop()` method as follows:

In [13]:

```
# in pandas axis=0 represents rows (default) and axis=1 represents columns.  
df_can.drop(['AREA', 'REG', 'DEV', 'Type', 'Coverage'], axis=1, inplace=True)  
df_can.head(2)
```

Out[13]:

	OdName	AreaName	RegName	DevName	1980	1981	1982	1983	1984	1985	...	2004	2005	2006
0	Afghanistan	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	...	2978	3436	3009
1	Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	...	1450	1223	856

2 rows × 38 columns



Let's rename the columns so that they make sense. We can use `rename()` method by passing in a dictionary of old and new names as follows:

In [14]:

```
df_can.rename(columns={'OdName':'Country', 'AreaName':'Continent', 'RegName':'Region'}, inplace=True)
df_can.columns
```

Out[14]:

```
Index([  'Country',  'Continent',  'Region',  'DevName',  1980,
        1981,      1982,      1983,      1984,      1985,
        1986,      1987,      1988,      1989,      1990,
        1991,      1992,      1993,      1994,      1995,
        1996,      1997,      1998,      1999,      2000,
        2001,      2002,      2003,      2004,      2005,
        2006,      2007,      2008,      2009,      2010,
        2011,      2012,      2013],
      dtype='object')
```

We will also add a 'Total' column that sums up the total immigrants by country over the entire period 1980 - 2013, as follows:

In [15]:

```
df_can['Total'] = df_can.sum(axis=1)
```

We can check to see how many null objects we have in the dataset as follows:

In [16]:

```
df_can.isnull().sum()
```

Out[16]:

Country	0
Continent	0
Region	0
DevName	0
1980	0
1981	0
1982	0
1983	0
1984	0
1985	0
1986	0
1987	0
1988	0
1989	0
1990	0
1991	0
1992	0
1993	0
1994	0
1995	0
1996	0
1997	0
1998	0
1999	0
2000	0
2001	0
2002	0
2003	0
2004	0
2005	0
2006	0
2007	0
2008	0
2009	0
2010	0
2011	0
2012	0
2013	0
Total	0

dtype: int64

Finally, let's view a quick summary of each column in our dataframe using the `describe()` method.


In [17]:

```
df_can.describe()
```

Out[17]:

	1980	1981	1982	1983	1984	1985	1986	
count	195.000000	195.000000	195.000000	195.000000	195.000000	195.000000	195.000000	1
mean	508.394872	566.989744	534.723077	387.435897	376.497436	358.861538	441.271795	6
std	1949.588546	2152.643752	1866.997511	1204.333597	1198.246371	1079.309600	1225.576630	21
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.500000	
50%	13.000000	10.000000	11.000000	12.000000	13.000000	17.000000	18.000000	
75%	251.500000	295.500000	275.000000	173.000000	181.000000	197.000000	254.000000	4
max	22045.000000	24796.000000	20620.000000	10015.000000	10170.000000	9564.000000	9470.000000	213

8 rows × 35 columns



pandas Intermediate: Indexing and Selection (slicing)

Select Column

There are two ways to filter on a column name:

Method 1: Quick and easy, but only works if the column name does NOT have spaces or special characters.

```
df.column_name # returns series
```

Method 2: More robust, and can filter on multiple columns.

```
df['column'] # returns series
df[['column 1', 'column 2']] # returns dataframe
```

Example: Let's try filtering on the list of countries ('Country').

In [18]:

```
df_can.Country # returns a series
```

Out[18]:

```
0      Afghanistan
1        Albania
2        Algeria
3  American Samoa
4        Andorra
...
190      Viet Nam
191  Western Sahara
192        Yemen
193        Zambia
194        Zimbabwe
Name: Country, Length: 195, dtype: object
```

Let's try filtering on the list of countries ('Country') and the data for years: 1980 - 1985.

In [19]:

```
df_can[['Country', 1980, 1981, 1982, 1983, 1984, 1985]] # returns a dataframe
# notice that 'Country' is string, and the years are integers.
# for the sake of consistency, we will convert all column names to string later on.
```

Out[19]:

	Country	1980	1981	1982	1983	1984	1985
0	Afghanistan	16	39	39	47	71	340
1	Albania	1	0	0	0	0	0
2	Algeria	80	67	71	69	63	44
3	American Samoa	0	1	0	0	0	0
4	Andorra	0	0	0	0	0	0
...
190	Viet Nam	1191	1829	2162	3404	7583	5907
191	Western Sahara	0	0	0	0	0	0
192	Yemen	1	2	1	6	0	18
193	Zambia	11	17	11	7	16	9
194	Zimbabwe	72	114	102	44	32	29

195 rows × 7 columns

Select Row

There are main 2 ways to select rows:

```
df.loc[label] # filters by the labels of the index/column
df.iloc[index] # filters by the positions of the index/column
```

Before we proceed, notice that the default index of the dataset is a numeric range from 0 to 194. This makes it very difficult to do a query by a specific country. For example to search for data on Japan, we need to know the corresponding index value.

This can be fixed very easily by setting the 'Country' column as the index using `set_index()` method.

In [20]:

```
df_can.set_index('Country', inplace=True)
# tip: The opposite of set is reset. So to reset the index, we can use df_can.reset_index()
```

In [21]:

```
df_can.head(3)
```

Out[21]:

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007
Country														
Afghanistan	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	496	...	3436	3009	264
Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	1	...	1223	856	70
Algeria	Africa	Northern Africa	Developing regions	80	67	71	69	63	44	69	...	3626	4807	362

3 rows × 38 columns



In [22]:

```
# optional: to remove the name of the index
df_can.index.name = None
```

Example: Let's view the number of immigrants from Japan (row 87) for the following scenarios: 1. The full row data (all columns) 2. For year 2013 3. For years 1980 to 1985

In [23]:

```
# 1. the full row data (all columns)
df_can.loc['Japan']
```

Out[23]:

Continent	Asia
Region	Eastern Asia
DevName	Developed regions
1980	701
1981	756
1982	598
1983	309
1984	246
1985	198
1986	248
1987	422
1988	324
1989	494
1990	379
1991	506
1992	605
1993	907
1994	956
1995	826
1996	994
1997	924
1998	897
1999	1083
2000	1010
2001	1092
2002	806
2003	817
2004	973
2005	1067
2006	1212
2007	1250
2008	1284
2009	1194
2010	1168
2011	1265
2012	1214
2013	982
Total	27707

Name: Japan, dtype: object

In [24]:

```
# alternate methods
df_can.iloc[87]
```

Out[24]:

Continent Asia
Region Eastern Asia
DevName Developed regions
1980 701
1981 756
1982 598
1983 309
1984 246
1985 198
1986 248
1987 422
1988 324
1989 494
1990 379
1991 506
1992 605
1993 907
1994 956
1995 826
1996 994
1997 924
1998 897
1999 1083
2000 1010
2001 1092
2002 806
2003 817
2004 973
2005 1067
2006 1212
2007 1250
2008 1284
2009 1194
2010 1168
2011 1265
2012 1214
2013 982
Total 27707
Name: Japan, dtype: object

In [25]:

```
df_can[df_can.index == 'Japan']
```

Out[25]:

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008
Japan	Asia	Eastern Asia	Developed regions	701	756	598	309	246	198	248	...	1067	1212	1250	1284

1 rows × 38 columns



In [26]:

```
# 2. for year 2013
df_can.loc['Japan', 2013]
```

Out[26]:

982

In [27]:

```
# alternate method
# year 2013 is the last column, with a positional index of 36
df_can.iloc[87, 36]
```

Out[27]:

982

In [28]:

```
# 3. for years 1980 to 1985
df_can.loc['Japan', [1980, 1981, 1982, 1983, 1984, 1984]]
```

Out[28]:

```
1980    701
1981    756
1982    598
1983    309
1984    246
1984    246
Name: Japan, dtype: object
```

In [29]:

```
# Alternative Method
df_can.iloc[87, [3, 4, 5, 6, 7, 8]]
```

Out[29]:

```
1980    701
1981    756
1982    598
1983    309
1984    246
1985    198
Name: Japan, dtype: object
```

Column names that are integers (such as the years) might introduce some confusion. For example, when we are referencing the year 2013, one might confuse that when the 2013th positional index.

To avoid this ambiguity, let's convert the column names into strings: '1980' to '2013'.

In [30]:

```
df_can.columns = list(map(str, df_can.columns))
# [print (type(x)) for x in df_can.columns.values] #<-- uncomment to check type of column headers
```

Since we converted the years to string, let's declare a variable that will allow us to easily call upon the full range of years:

In [31]:

```
# useful for plotting later on
years = list(map(str, range(1980, 2014)))
years
```

Out[31]:

```
['1980',
 '1981',
 '1982',
 '1983',
 '1984',
 '1985',
 '1986',
 '1987',
 '1988',
 '1989',
 '1990',
 '1991',
 '1992',
 '1993',
 '1994',
 '1995',
 '1996',
 '1997',
 '1998',
 '1999',
 '2000',
 '2001',
 '2002',
 '2003',
 '2004',
 '2005',
 '2006',
 '2007',
 '2008',
 '2009',
 '2010',
 '2011',
 '2012',
 '2013']
```

Filtering based on a criteria

To filter the dataframe based on a condition, we simply pass the condition as a boolean vector.

For example, Let's filter the dataframe to show the data on Asian countries (AreaName = Asia).

In [32]:

```
# 1. create the condition boolean series
condition = df_can['Continent'] == 'Asia'
print(condition)
```

Afghanistan	True
Albania	False
Algeria	False
American Samoa	False
Andorra	False
...	
Viet Nam	True
Western Sahara	False
Yemen	True
Zambia	False
Zimbabwe	False

Name: Continent, Length: 195, dtype: bool

In [33]:


```
# 2. pass this condition into the dataframe  
df_can[condition]
```


	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006
Afghanistan	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	496	...	3436	3009
Armenia	Asia	Western Asia	Developing regions	0	0	0	0	0	0	0	...	224	218
Azerbaijan	Asia	Western Asia	Developing regions	0	0	0	0	0	0	0	...	359	236
Bahrain	Asia	Western Asia	Developing regions	0	2	1	1	1	3	0	...	12	12
Bangladesh	Asia	Southern Asia	Developing regions	83	84	86	81	98	92	486	...	4171	4014
Bhutan	Asia	Southern Asia	Developing regions	0	0	0	0	1	0	0	...	5	10
Brunei Darussalam	Asia	South-Eastern Asia	Developing regions	79	6	8	2	2	4	12	...	4	5
Cambodia	Asia	South-Eastern Asia	Developing regions	12	19	26	33	10	7	8	...	370	529
China	Asia	Eastern Asia	Developing regions	5123	6682	3308	1863	1527	1816	1960	...	42584	33518
China, Hong Kong Special Administrative Region	Asia	Eastern Asia	Developing regions	0	0	0	0	0	0	0	...	729	712
China, Macao Special Administrative Region	Asia	Eastern Asia	Developing regions	0	0	0	0	0	0	0	...	21	32
Cyprus	Asia	Western Asia	Developing regions	132	128	84	46	46	43	48	...	7	9
Democratic People's Republic of Korea	Asia	Eastern Asia	Developing regions	1	1	3	1	4	3	0	...	14	10
Georgia	Asia	Western Asia	Developing regions	0	0	0	0	0	0	0	...	114	125
India	Asia	Southern Asia	Developing regions	8880	8670	8147	7338	5704	4211	7150	...	36210	33848
Indonesia	Asia	South-Eastern Asia	Developing regions	186	178	252	115	123	100	127	...	632	613
Iran (Islamic Republic of)	Asia	Southern Asia	Developing regions	1172	1429	1822	1592	1977	1648	1794	...	5837	7480
Iraq	Asia	Western Asia	Developing regions	262	245	260	380	428	231	265	...	2226	1788
Israel	Asia	Western Asia	Developing regions	1403	1711	1334	541	446	680	1212	...	2446	2625
Japan	Asia	Eastern Asia	Developed regions	701	756	598	309	246	198	248	...	1067	1212
Jordan	Asia	Western Asia	Developing regions	177	160	155	113	102	179	181	...	1940	1827

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006
Kazakhstan	Asia	Central Asia	Developing regions	0	0	0	0	0	0	0	...	506	408
Kuwait	Asia	Western Asia	Developing regions	1	0	8	2	1	4	4	...	66	35
Kyrgyzstan	Asia	Central Asia	Developing regions	0	0	0	0	0	0	0	...	173	161
Lao People's Democratic Republic	Asia	South-Eastern Asia	Developing regions	11	6	16	16	7	17	21	...	42	74
Lebanon	Asia	Western Asia	Developing regions	1409	1119	1159	789	1253	1683	2576	...	3709	3802
Malaysia	Asia	South-Eastern Asia	Developing regions	786	816	813	448	384	374	425	...	593	580
Maldives	Asia	Southern Asia	Developing regions	0	0	0	1	0	0	0	...	0	0
Mongolia	Asia	Eastern Asia	Developing regions	0	0	0	0	0	0	0	...	59	64
Myanmar	Asia	South-Eastern Asia	Developing regions	80	62	46	31	41	23	18	...	210	953
Nepal	Asia	Southern Asia	Developing regions	1	1	6	1	2	4	13	...	607	540
Oman	Asia	Western Asia	Developing regions	0	0	0	8	0	0	0	...	14	18
Pakistan	Asia	Southern Asia	Developing regions	978	972	1201	900	668	514	691	...	14314	13127
Philippines	Asia	South-Eastern Asia	Developing regions	6051	5921	5249	4562	3801	3150	4166	...	18139	18400
Qatar	Asia	Western Asia	Developing regions	0	0	0	0	0	0	1	...	11	2
Republic of Korea	Asia	Eastern Asia	Developing regions	1011	1456	1572	1081	847	962	1208	...	5832	6215
Saudi Arabia	Asia	Western Asia	Developing regions	0	0	1	4	1	2	5	...	198	252
Singapore	Asia	South-Eastern Asia	Developing regions	241	301	337	169	128	139	205	...	392	298
Sri Lanka	Asia	Southern Asia	Developing regions	185	371	290	197	1086	845	1838	...	4930	4714
State of Palestine	Asia	Western Asia	Developing regions	0	0	0	0	0	0	0	...	453	627
Syrian Arab Republic	Asia	Western Asia	Developing regions	315	419	409	269	264	385	493	...	1458	1145
Tajikistan	Asia	Central Asia	Developing regions	0	0	0	0	0	0	0	...	85	46
Thailand	Asia	South-Eastern Asia	Developing regions	56	53	113	65	82	66	78	...	575	500
Turkey	Asia	Western Asia	Developing regions	481	874	706	280	338	202	257	...	2065	1638

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006
Turkmenistan	Asia	Central Asia	Developing regions	0	0	0	0	0	0	0	...	40	26
United Arab Emirates	Asia	Western Asia	Developing regions	0	2	2	1	2	0	5	...	31	42
Uzbekistan	Asia	Central Asia	Developing regions	0	0	0	0	0	0	0	...	330	262
Viet Nam	Asia	South-Eastern Asia	Developing regions	1191	1829	2162	3404	7583	5907	2741	...	1852	3153
Yemen	Asia	Western Asia	Developing regions	1	2	1	6	0	18	7	...	161	140

49 rows × 38 columns



In [34]:

```
# we can pass multiple criteria in the same line.
# Let's filter for AreaName = Asia and RegName = Southern Asia

df_can[(df_can['Continent']=='Asia') & (df_can['Region']=='Southern Asia')]

# note: When using 'and' and 'or' operators, pandas requires we use '&' and '|' instead of 'and' and 'or'
# don't forget to enclose the two conditions in parentheses
```

Out[34]:

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006
Afghanistan	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	496	...	3436	3009
Bangladesh	Asia	Southern Asia	Developing regions	83	84	86	81	98	92	486	...	4171	4014
Bhutan	Asia	Southern Asia	Developing regions	0	0	0	0	1	0	0	...	5	10
India	Asia	Southern Asia	Developing regions	8880	8670	8147	7338	5704	4211	7150	...	36210	33848
Iran (Islamic Republic of)	Asia	Southern Asia	Developing regions	1172	1429	1822	1592	1977	1648	1794	...	5837	7480
Maldives	Asia	Southern Asia	Developing regions	0	0	0	1	0	0	0	...	0	0
Nepal	Asia	Southern Asia	Developing regions	1	1	6	1	2	4	13	...	607	540
Pakistan	Asia	Southern Asia	Developing regions	978	972	1201	900	668	514	691	...	14314	13127
Sri Lanka	Asia	Southern Asia	Developing regions	185	371	290	197	1086	845	1838	...	4930	4714

9 rows × 38 columns



Before we proceed: let's review the changes we have made to our dataframe.

In [35]:

```
print('data dimensions:', df_can.shape)
print(df_can.columns)
df_can.head(2)
```

```
data dimensions: (195, 38)
Index(['Continent', 'Region', 'DevName', '1980', '1981', '1982', '1983',
      '1984', '1985', '1986', '1987', '1988', '1989', '1990', '1991', '1992',
      '1993', '1994', '1995', '1996', '1997', '1998', '1999', '2000', '2001',
      '2002', '2003', '2004', '2005', '2006', '2007', '2008', '2009', '2010',
      '2011', '2012', '2013', 'Total'],
      dtype='object')
```

Out[35]:

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007
Afghanistan	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	496	...	3436	3009	2600
Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	1	...	1223	856	700

2 rows × 38 columns

Visualizing Data using Matplotlib

Matplotlib: Standard Python Visualization Library

The primary plotting library we will explore in the course is [Matplotlib \(http://matplotlib.org/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01\)](http://matplotlib.org/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01). As mentioned on their website:

Matplotlib is a Python 2D plotting library which produces publication quality figures in a variety of hardcopy formats and interactive environments across platforms. Matplotlib can be used in Python scripts, the Python and IPython shell, the jupyter notebook, web application servers, and four graphical user interface toolkits.

If you are aspiring to create impactful visualization with python, Matplotlib is an essential tool to have at your disposal.

Matplotlib.Pyplot

One of the core aspects of Matplotlib is `matplotlib.pyplot`. It is Matplotlib's scripting layer which we studied in details in the videos about Matplotlib. Recall that it is a collection of command style functions that make Matplotlib work like MATLAB. Each `pyplot` function makes some change to a figure: e.g., creates a figure, creates a plotting area in a figure, plots some lines in a plotting area, decorates the plot with labels, etc. In this lab, we will work with the scripting layer to learn how to generate line plots. In future labs, we will get to work with the Artist layer as well to experiment first hand how it differs from the scripting layer.

Let's start by importing `matplotlib` and `matplotlib.pyplot` as follows:

In [36]:

```
# we are using the inline backend
%matplotlib inline

import matplotlib as mpl
import matplotlib.pyplot as plt
```

*optional: check if Matplotlib is loaded.

In [37]:

```
print('Matplotlib version: ', mpl.__version__) # >= 2.0.0
```

Matplotlib version: 3.3.4

*optional: apply a style to Matplotlib.

In [38]:

```
print(plt.style.available)
mpl.style.use(['ggplot']) # optional: for ggplot-like style
```

```
['Solarize_Light2', '_classic_test_patch', 'bmh', 'classic', 'dark_background', 'fast',
'fivethirtyeight', 'ggplot', 'grayscale', 'seaborn', 'seaborn-bright', 'seaborn-colorblind',
'seaborn-dark', 'seaborn-dark-palette', 'seaborn-darkgrid', 'seaborn-deep', 'seaborn-muted',
'seaborn-notebook', 'seaborn-paper', 'seaborn-pastel', 'seaborn-poster', 'seaborn-talk',
'seaborn-ticks', 'seaborn-white', 'seaborn-whitegrid', 'tableau-colorblind10']
```

Plotting in *pandas*

Fortunately, *pandas* has a built-in implementation of Matplotlib that we can use. Plotting in *pandas* is as simple as appending a `.plot()` method to a series or dataframe.

Documentation:

- [Plotting with Series \(http://pandas.pydata.org/pandas-docs/stable/api.html?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01#plotting\)](http://pandas.pydata.org/pandas-docs/stable/api.html?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01#plotting)
- [Plotting with Dataframes \(http://pandas.pydata.org/pandas-docs/stable/api.html?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01#api-dataframe-plotting\)](http://pandas.pydata.org/pandas-docs/stable/api.html?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDeveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01#api-dataframe-plotting)

Line Plots (Series/Dataframe)

What is a line plot and why use it?

A line chart or line plot is a type of plot which displays information as a series of data points called 'markers' connected by straight line segments. It is a basic type of chart common in many fields. Use line plot when you have a continuous data set. These are best suited for trend-based visualizations of data over a period of time.

Let's start with a case study:

In 2010, Haiti suffered a catastrophic magnitude 7.0 earthquake. The quake caused widespread devastation and loss of life and about three million people were affected by this natural disaster. As part of Canada's humanitarian effort, the Government of Canada stepped up its effort in accepting refugees from Haiti. We can quickly visualize this effort using a Line plot:

Question: Plot a line graph of immigration from Haiti using `df.plot()` .

First, we will extract the data series for Haiti.

In [39]:

```
haiti = df_can.loc['Haiti', years] # passing in years 1980 - 2013 to exclude the 'total' column
haiti.head()
```

Out[39]:

```
1980    1666
1981    3692
1982    3498
1983    2860
1984    1418
Name: Haiti, dtype: object
```


Next, we will plot a line plot by appending `.plot()` to the `haiti` dataframe.

In [40]:

```
haiti.plot()
```

Out[40]:

<AxesSubplot:>



`pandas` automatically populated the x-axis with the index values (years), and the y-axis with the column values (population). However, notice how the years were not displayed because they are of type `string`. Therefore, let's change the type of the index values to `integer` for plotting.


Also, let's label the x and y axis using `plt.title()` , `plt.ylabel()` , and `plt.xlabel()` as follows:

In [41]:

```
haiti.index = haiti.index.map(int) # Let's change the index values of Haiti to type integer for plotting
haiti.plot(kind='line')

plt.title('Immigration from Haiti')
plt.ylabel('Number of immigrants')
plt.xlabel('Years')

plt.show() # need this line to show the updates made to the figure
```



We can clearly notice how number of immigrants from Haiti spiked up from 2010 as Canada stepped up its efforts to accept refugees from Haiti. Let's annotate this spike in the plot by using the `plt.text()` method.


In [42]:

```
haiti.plot(kind='line')

plt.title('Immigration from Haiti')
plt.ylabel('Number of Immigrants')
plt.xlabel('Years')

# annotate the 2010 Earthquake.
# syntax: plt.text(x, y, label)
plt.text(2000, 6000, '2010 Earthquake') # see note below

plt.show()
```



With just a few lines of code, you were able to quickly identify and visualize the spike in immigration!

Quick note on x and y values in `plt.text(x, y, label)` :

Since the x-axis (years) is type 'integer', we specified x as a year. The y axis (number of immigrants) is type 'integer', so we can just specify the value y = 6000.

```
plt.text(2000, 6000, '2010 Earthquake') # years stored as type int
```

If the years were stored as type 'string', we would need to specify x as the index position of the year. Eg 20th index is year 2000 since it is the 20th year with a base year of 1980.

```
plt.text(20, 6000, '2010 Earthquake') # years stored as type int
```

We will cover advanced annotation methods in later modules.

We can easily add more countries to line plot to make meaningful comparisons immigration from different countries.

Question: Let's compare the number of immigrants from India and China from 1980 to 2013.

Step 1: Get the data set for China and India, and display the dataframe.

In [47]:

```
### type your answer here
# df_CI = ...
# df_CI
df_CI = df_can.loc[['India', 'China'], years]
df_CI
```

Out[47]:

	1980	1981	1982	1983	1984	1985	1986	1987	1988	1989	...	2004	2005	2006	2007	2008
India	8880	8670	8147	7338	5704	4211	7150	10189	11522	10343	...	28235	36210	33848	28742	28121
China	5123	6682	3308	1863	1527	1816	1960	2643	2758	4323	...	36619	42584	33518	27642	30483

2 rows × 34 columns



[Click here for a sample python solution](#)

Step 2: Plot graph. We will explicitly specify line plot by passing in `kind` parameter to `plot()` .


In [48]:

```
df_CI.plot(kind='line')
```

Out[48]:

<AxesSubplot:>

1980
1981
1982
1983
1984
1985
1986
1987
1988
1989
1990
1991
1992
1993
1994
1995
1996
1997
1998
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2009
2010
2011
2012
2013



[Click here for a sample python solution](#)

That doesn't look right...

Recall that *pandas* plots the indices on the x-axis and the columns as individual lines on the y-axis. Since `df_CI` is a dataframe with the `country` as the index and `years` as the columns, we must first transpose the dataframe using `transpose()` method to swap the row and columns.

In [49]:

```
df_CI = df_CI.transpose()
df_CI.head()
```

Out[49]:

	India	China
1980	8880	5123
1981	8670	6682
1982	8147	3308
1983	7338	1863
1984	5704	1527

pandas will auomatically graph the two countries on the same graph. Go ahead and plot the new transposed dataframe. Make sure to add a title to the plot and label the axes.


In [50]:

```
### type your answer here
```

```
df_CI.index = df_CI.index.map(int) # let's change the index values of df_CI to type integer for plotting
df_CI.plot(kind='line')

plt.title('Immigrants from China and India')
plt.ylabel('Number of Immigrants')
plt.xlabel('Years')

plt.show()
```



[Click here for a sample python solution](#)

From the above plot, we can observe that the China and India have very similar immigration trends through the years.

Note: How come we didn't need to transpose Haiti's dataframe before plotting (like we did for df_CI)?

That's because `haiti` is a series as opposed to a dataframe, and has the years as its indices as shown below.

```
print(type(haiti))  
print(haiti.head(5))
```

```
class 'pandas.core.series.Series'  
1980 1666  
1981 3692  
1982 3498  
1983 2860  
1984 1418  
Name: Haiti, dtype: int64
```

Line plot is a handy tool to display several dependent variables against one independent variable. However, it is recommended that no more than 5-10 lines on a single graph; any more than that and it becomes difficult to interpret.

Question: Compare the trend of top 5 countries that contributed the most to immigration to Canada.

In [51]:

```
### type your answer here
```

```
#The correct answer is:
```

```
#Step 1: Get the dataset. Recall that we created a Total column that calculates cumulative immigration by country.
```

```
#We will sort on this column to get our top 5 countries using pandas sort_values() method.
```

```
inplace = True # parameter saves the changes to the original df_can dataframe  
df_can.sort_values(by='Total', ascending=False, axis=0, inplace=True)
```

```
# get the top 5 entries
```

```
df_top5 = df_can.head(5)
```

```
# transpose the dataframe
```

```
df_top5 = df_top5[years].transpose()
```

```
print(df_top5)
```

```
#Step 2: Plot the dataframe. To make the plot more readable, we will change the size using the `figsize` parameter.
```

```
df_top5.index = df_top5.index.map(int) # Let's change the index values of df_top5 to type integer for plotting
```

```
df_top5.plot(kind='line', figsize=(14, 8)) # pass a tuple (x, y) size
```

```
plt.title('Immigration Trend of Top 5 Countries')
```

```
plt.ylabel('Number of Immigrants')
```


```
plt.xlabel('Years')
```

```
plt.show()
```


	India	China	United Kingdom of Great Britain and Northern Ireland	\
1980	8880	5123		22045
1981	8670	6682		24796
1982	8147	3308		20620
1983	7338	1863		10015
1984	5704	1527		10170
1985	4211	1816		9564
1986	7150	1960		9470
1987	10189	2643		21337
1988	11522	2758		27359
1989	10343	4323		23795
1990	12041	8076		31668
1991	13734	14255		23380
1992	13673	10846		34123
1993	21496	9817		33720
1994	18620	13128		39231
1995	18489	14398		30145
1996	23859	19415		29322
1997	22268	20475		22965
1998	17241	21049		10367
1999	18974	30069		7045
2000	28572	35529		8840
2001	31223	36434		11728
2002	31889	31961		8046
2003	27155	36439		6797
2004	28235	36619		7533
2005	36210	42584		7258
2006	33848	33518		7140
2007	28742	27642		8216
2008	28261	30037		8979
2009	29456	29622		8876
2010	34235	30391		8724
2011	27509	28502		6204
2012	30933	33024		6195
2013	33087	34129		5827

	Philippines	Pakistan
1980	6051	978
1981	5921	972
1982	5249	1201
1983	4562	900
1984	3801	668
1985	3150	514
1986	4166	691
1987	7360	1072
1988	8639	1334
1989	11865	2261
1990	12509	2470
1991	12718	3079
1992	13670	4071
1993	20479	4777
1994	19532	4666
1995	15864	4994
1996	13692	9125
1997	11549	13073
1998	8735	9068
1999	9734	9979
2000	10763	15400
2001	13836	16708
2002	11707	15110
2003	12758	13205

2004	14004	13399
2005	18139	14314
2006	18400	13127
2007	19837	10124
2008	24887	8994
2009	28573	7217
2010	38617	6811
2011	36765	7468
2012	34315	11227
2013	29544	12603



[Click here for a sample python solution](#)

Other Plots

Congratulations! you have learned how to wrangle data with python and create a line plot with Matplotlib. There are many other plotting styles available other than the default Line plot, all of which can be accessed by passing `kind` keyword to `plot()` . The full list of available plots are as follows:

- `bar` for vertical bar plots
- `barh` for horizontal bar plots
- `hist` for histogram
- `box` for boxplot
- `kde` or `density` for density plots
- `area` for area plots
- `pie` for pie plots
- `scatter` for scatter plots
- `hexbin` for hexbin plot

Thank you for completing this lab!

Author

[Alex Aklson \(https://www.linkedin.com/in/aklson/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01\)](https://www.linkedin.com/in/aklson/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01)

Other Contributors

[Jay Rajasekharan \(https://www.linkedin.com/in/jayrajasekharan/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01\)](https://www.linkedin.com/in/jayrajasekharan/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01),
[Ehsan M. Kermani \(https://www.linkedin.com/in/ehsanmkermani/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01\)](https://www.linkedin.com/in/ehsanmkermani/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01),
[Slobodan Markovic \(https://www.linkedin.com/in/slobodan-markovic/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01\)](https://www.linkedin.com/in/slobodan-markovic/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01),
[Weiqing Wang \(https://www.linkedin.com/in/weiqing-wang-641640133/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01\)](https://www.linkedin.com/in/weiqing-wang-641640133/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=10006555&utm_id=NA-SkillsNetwork-Channel-SkillsNetworkCoursesIBMDveloperSkillsNetworkDV0101ENSkillsNetwork20297740-2021-01-01)

Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2021-05-29	2.4	Weiqing Wang	Fixed typos and code smells.
2021-01-20	2.3	Lakshmi Holla	Changed TOC cell markdown
2020-11-20	2.2	Lakshmi Holla	Changed IBM box URL
2020-11-03	2.1	Lakshmi Holla	Changed URL and info method
2020-08-27	2.0	Lavanya	Moved Lab to course repo in GitLab

© IBM Corporation 2020. All rights reserved.

In []: